Claude Code — Configuration Report

Prepared for **Murpheux** · 28 September 2026 · Verified against **Claude Code 2.1.283** (latest in `anthropics/claude-code` CHANGELOG)
Folder: `~/agent-configs/claude-code/` → installs to `~/.claude/`

1. Summary

Claude Code is the closest match to the OpenCode design. Since 2026 its `settings.json` accepts `"agent": "router"`, so the **main session itself runs as a cheap Haiku router**. The shared guard hook keeps that main thread read-only, which forces delegation to seven specialist subagents on the right model tier. Enforcement is layered: `permissions.deny` (always wins, even in bypass mode), the PreToolUse guard hook, the OS sandbox, and credential scrubbing of subprocess environments.

8 agents	**16** slash commands	**20** skills	**137** allow · **9** ask · **66** deny	**5** MCP servers

2. Sources used

- **hesreallyhim/awesome-claude-code** (README: agent skills, workflows, hooks, status lines, usage monitors, orchestrators).
- Official: `anthropics/claude-code` CHANGELOG + plugins, `anthropics/skills`, `code.claude.com/docs` (settings, permissions, sub-agents, skills, hooks, MCP, sandboxing, model-config, env-vars), `agentskills.io` spec.

3. What was built

File	Purpose
`.claude/settings.json`	`agent: router`, model `sonnet` + effort medium for direct sessions, `autoCompactWindow 200000`, permissions (allow read-only/validation, ask for commits/plan/apply/helm/op, deny destructive + secret files), guard hook, sandbox, env, status line
`.claude/settings.gateway.json`	optional overlay (`./install.sh --gateway`): `ANTHROPIC_BASE_URL=https://scarlet:4000`, key via `apiKeyHelper` (`op read` executed by Claude Code, never the model), gateway model discovery, `NODE_EXTRA_CA_CERTS`
`.claude/CLAUDE.md`	full global instructions (every fact of the original AGENTS.md; Claude-specific §13/§14/§17)
`.claude/agents/*.md`	router, explore (overrides built-in), reviewer, planner, developer, sre, architect, security
`.claude/commands/*.md`	16 commands with `!`…`` context injection, pre-approved read-only `allowed-tools`, `model: haiku` on cheap ones
`.claude/skills/`	20 skills (name == folder, ≤1024-char descriptions)
`.claude/hooks/guard.py` + `guardrails.json`	shared guard with `--main-thread-readonly` router gate

`.claude/statusline.py`	model · effort · agent · $ cost · context % (warns ≥ 70%)
`mcp/add-mcp-servers.sh`	`claude mcp add-json --scope user`: github (read-only, 5 toolsets), context7, mslearn, awsdocs, terraform — secrets as runtime `${VAR}`

4. Agent tiers

Agent	Model · effort	Tools / notes
router (main session)	haiku	read/search/web + Agent; Edit/Write and non-read-only Bash blocked by the hook
Explore (override)	haiku	read-only; the built-in inherits the main model since 2.1.198
reviewer	haiku	read + git; `omitClaudeMd` keeps its context small
planner	opus · medium	`permissionMode: plan` (read-only)
developer / sre	sonnet · medium	all tools / read + shell
architect	opus · high	read + docs/diagram writes
security	opus · high	read + scanners (premium, HIGH_RISK only)

5. Cost levers

- Haiku router plus isolated subagent contexts: the router never carries implementation transcripts.
- Explore overridden to Haiku; `/commit` `/explain` `/summarize` `/pr` pinned to Haiku.
- `autoCompactWindow: 200000`, `BASH_MAX_OUTPUT_LENGTH=20000`, subagent depth 2 / 4 concurrent, `maxTurns` per agent.
- Skills load on demand; a status line shows $ and context % live; `npx ccusage` for history.
- Not used: `advisorModel`. Each advisor call re-reads the whole conversation uncached, and it needs pass-through to the Anthropic API.

6. Guardrails

1. **permissions.deny**: push, reset --hard, filter-branch, root/home rm, disk tools, destroy, `-auto-approve`, bucket/namespace deletion, secret-printing commands, secret files, `.git/**` edits, 1Password MCP.
2. **Guard hook**: 33 regex rules and 18 protected-path patterns, for cases prefix rules miss (`git -C x push -f`, `curl | sh`, `.env.*` with example exceptions, interactive programs). It also runs the router gate.
3. **Sandbox**: writes confined to the project; network limited to package registries and docs, and a new domain prompts; cloud and cluster CLIs are excluded and go through normal prompts.
4. `CLAUDE_CODE_SUBPROCESS_ENV_SCRUB=1`, `DISABLE_TELEMETRY=1`, `enableAllProjectMcpServers: false`.

7. From awesome-claude-code

Entry	Decision
dshakes/compass (cost-tiered agents + router)	Idea adopted natively; its installer (replaces `~/.claude` with symlinks) not used
kenryu42/claude-code-safety-net	Idea adopted → `guard.py`
sirmalloc/ccstatusline, claude-hud	Idea adopted → `statusline.py`
ccusage	Install (cost reports)
trailofbits/skills, antonbabenko/terraform-skill	Install (security and Terraform skills)

jgraph/drawio-mcp plugin, Anthropic security-guidance plugin	Install (listed in `add-mcp-servers.sh`)
dagger/container-use	Install for isolated parallel/bypass work
ypollak2/llm-router	Skip; LiteLLM already routes, so only the policy was copied

8. Gotchas handled

- Deny beats allow and can't be carved out, so `.env.example` exceptions live in the hook rather than in deny rules.
- Bash rules are text-based and can be routed around (`sh -c`, absolute paths); the hook and sandbox are the real boundary.
- A leading `/` in user settings means `~/.claude/`, so all path rules use `~/` or `**/`.
- Hooks fail open on a wrong path or a timeout; the installer runs the guard self-test and aborts on failure.
- `CLAUDE_CODE_SUBAGENT_MODEL` no longer overrides frontmatter, so every agent pins its own model.
- A non-Anthropic gateway disables MCP tool search, which is why the gateway is an opt-in overlay.

9. Validation

JSON/frontmatter/skill-spec validation passed. The guard's 70 self-tests passed. The installer smoke test in a temp HOME confirmed: 8 agents, 16 commands and 20 skills installed; the hook denies a force-push from a subagent; the router gate denies a main-thread Edit; the status line renders. The instructions retain all 277 checked items of the original AGENTS.md. Claude Code isn't installed on this workstation, so there was no live run.

10. Install & next steps

```
cd ~/agent-configs/claude-code && DRY_RUN=1 ./install.sh && ./install.sh      # add --gateway for Agentgateway
./mcp/add-mcp-servers.sh
claude            # router session;  direct session: ROUTER_MODE=off claude --agent developer
# WSL2 sandbox: sudo apt install bubblewrap socat
```